

```
figsize( 9, 3.5 )
np.set_printoptions(precision=3, suppress= True)
import scipy.stats as stats
```

Chapter 2: Modeling with PyMC

This chapter introduces more PyMC syntax and design patterns, and ways to think about Bayesian modeling a system.

A little more on PyMC

Parent and Child relationships

To assist with terminology, and to be consistent with PyMC's documentation, we introduce *parent* and *children* variables.

- *parent variables* are variables that influence another variable.
- *children variable* are variables that are affected by other variables, i.e. are the subject of parent variables.

Variables can be both parent and children variables. For example, consider the PyMC code below

```
import pymc as mc

poi_parameter = mc.Exponential( "Poisson_param", 1 )

data_generator = mc.Poisson("data_generator", poi_parameter )
```

`poi_parameter` controls the parameter of `data_generator`, hence influences its value. The former is a parent of the latter. By symmetry, `data_generator` is a child of `poi_parameter`.

This nomenclature is introduced to help us describe relationships in PyMC modeling. You can access a variables children and parent variables using the `children` and `parents` methods attached to variables.

```
print poi_parameter.children
print
print data_generator.parents

set([<pymc.distributions.Poisson 'data_generator' at 0x05847E10>])

{'mu': <pymc.distributions.Exponential 'Poisson_param' at 0x05847CF0>}
```

Of course a child can have more than one parent, and parent can have many children.

PyMC Variables

All PyMC variables also expose a `value` attribute. This method produces the *current* (possible random) value of the variable, given the variable's parents. To use the same variables from before:

```
print poi_parameter.value
print data_generator.value
```

```
0.746568979507
```

```
1
```

PyMC is concerned with two types of programming variables: *stochastic* and *deterministic*.

- *stochastic variables* are variables that are not deterministic, i.e., even if you knew all the values of the variables' parents (if it even had any parents), it would still be random. Included in this category are instances of classes `Poisson`, `DiscreteUniform`, and `Exponential`.
- *deterministic variables* are variables that not random if the variables' parents are not random. This might be confusing at first: a quick mental check is *if I knew all of variable x 's parent variables, I could determine completely what x is*. This is relevant as we are dealing with random variables.

Stochastic variables

Initializing a stochastic variable requires a `name` argument, plus additional parameters that a class specific. For example:

```
some_variable = mc.DiscreteUniform( "discrete_uni_var", 0, 4 )
```

where 0,4 are the `DiscreteUniform`-specific bounds on the random variable. The [PyMC docs](#) contain the specific parameters for stochastic variables. (Or use ?? if you are using IPython!)

Rather than creating a Python array of stochastic variables, addressing the `size` keyword in the call to `Stochastic` creates multivariate array of the (independent) stochastic variables. The array behaves like a Numpy array when used like one, and references to its `value` attribute return Numpy arrays.

We can also call on a stochastic variables `random()` method, which (given the parent values) will generate a new (random) value.

```
lambda_1 = mc.Exponential( "lambda_1", 1 )
lambda_2 = mc.Exponential( "lambda_2", 1 )
tau = mc.DiscreteUniform( "tau", lower = 0, upper = 10 )
print "lambda_1.value = %.3f"%lambda_1.value
print "lambda_2.value = %.3f"%lambda_2.value
print "tau.value = %.3f"%tau.value
print

lambda_1.random(), lambda_2.random(), tau.random()
print "After calling random() on the variables..."
```

```
print "lambda_1.value = %.3f"%lambda_1.value
print "lambda_2.value = %.3f"%lambda_2.value
print "tau.value = %.3f"%tau.value
```

```
lambda_1.value = 1.160
lambda_2.value = 0.660
tau.value = 10.000
```

After calling `random()` on the variables...

```
lambda_1.value = 0.093
lambda_2.value = 0.236
tau.value = 1.000
```

The call to `random` stores a new value into the variables `value` attribute. In fact, this new value is stored in the computers cache for faster recall and efficiency.

Deterministic variables

Since most variables you will be modeling are stochastic, we distinguish deterministic variables with the `pymc.deterministic` wrapper. This is the easiest way, but not the only way, to create deterministic variables. This is not completely true: elementary operations, like addition, exponentials etc. implicitly create deterministic variables.

```
type( lambda_1 + lambda_2 )
```

```
pymc.PyMCObjects.Deterministic
```

The use of the `deterministic` wrapper was seen in the previous text-message example. Recall the model for λ looked like:

$$\lambda = \begin{cases} \lambda_1 & \text{if } t < \tau \\ \lambda_2 & \text{if } t \geq \tau \end{cases}$$

And in PyMC code:

```
@mc.deterministic
def lambda_( tau = tau, lambda_1 = lambda_1, lambda_2 = lambda_2 ):
    out = np.zeros( 10 )
    out[:tau] = lambda_1 #lambda before tau is lambda1
    out[tau:] = lambda_2 #lambda after tau is lambda1
    return out
```

Clearly, if τ , λ_1 and λ_2 are known, then λ is known completely, hence it is a deterministic variable.


Including observations

At this point, it may not look like it, but we have fully specified our priors. For example, we can ask

and answer questions like "What does my prior distribution of λ_1 look like?"

```
samples = [ lambda_1.random() for i in range( 20000) ]
hist( samples, bins = 35, normed=True )
plt.title( "Prior distribution for  $\lambda_1$ " )
```

<matplotlib.text.Text at 0x158382b0>



To frame this in the notation of the first chapter, though this is a slight abuse of notation, we have specified $P(A)$. Our next goal is to include data/evidence/observations X into our model.

PyMC stochastic variables have a keyword argument `observed` which accepts a boolean (False by default). `observed` has a very simple role: fix the variables current value, i.e. make `value` immutable. For this to make sense, we have to specify an initial `value` in the class call. For example:

```
fixed_variable = mc.Poisson( "fxd", 1, value = 10, observed = True )
print "value: ",fixed_variable.value
print "calling .random()"
fixed_variable.random()
print "value: ",fixed_variable.value
```

```
value: 10
calling .random()
value: 10
```

This is how we include data into our models: initializing the stochastic variable to have a *fixed value*.

...the stochastic variable to have a *fixed value*.

That might seem strange to hear at first. In fact, Bayesian analysis sees observed data as simply fixed parameters. Taking this to its logical conclusion, any predictions made by Bayesian analysis is seen as fitting another parameter in the model.

Modeling approaches

A good starting place to modeling is to think about *how your data might have been generated*.

Think from a god-like position, and try to think about how *you* would recreate the data. For example, in the last chapter, we investigated text message data:

- I. We started by thinking "what is the best random variable to describe this data?" A Poisson random variable is a good candidate because it assigns probabilities to count data.
- II. Next, we think, "Ok, assuming texts are Poisson-distributed, what do I need for the Poisson distribution?" Well, the Poisson distribution has a parameters λ .
- III. Do we know λ ? No. In fact, we have a suspicion that there are *two* λ values, one for the earlier behaviour and one for the latter behaviour. We don't know when the behaviour switches though.
- IV. What is a good distribution for the two λ s? The exponential is good, as it assigns probabilities to positive real numbers. Well the exponential distribution has a parameter too, call it α .
- V. Do we know what the parameter α might be? No. We could continue and assign a distribution to α , but it's better to stop once we reach a set level of ignorance: whereas we have a prior belief about λ , ("it probably changes over time", "it's likely between 1 and 8", etc.), we don't really have any strong beliefs about α . So it's best to stop here.
- VI. What is a good value for α then? We think that the λ s are between 1-8, so if we set α really low (which corresponds to larger probability on high values) we are not reflecting our prior well. Similar, a too-high α misses our prior as well. A good idea for α as to reflect our belief is to set the value so that the mean of λ , given α , is equal to our observed mean. This was shown in the last chapter.

In this way, we are thinking about how the data might have been created.

Example: Poisson Regression

Perhaps the most important result from medical research was the *now obvious* link between *smoking and cancer*. We'll try to establish a link using Bayesian methods. We have a decision here: should we include a prior that biases us towards there existing a significant link between smoking and cancer? I think we should act like scientists at the turn of the century, and assume there's is no *a priori* reason to assume a link.

The dataset we will be using contains 36 cohorts (a unique group with a larger population), each cohort has variables:

- age: in five-year age groups coded 1 to 9 for 40-44, 45-49, 50-54, 55-59, 60-64, 65-69, 70-74, 75-79, 80+.
- cigar/pipe smoking status: 1 if the cohort only smokes cigars/pipes, 0 else.
- cigar & cigarette smoking status: 1 if the cohort smoke *both* cigars/pipes and cigarettes, 0 else.
- cigarette smoking status: 1 if the cohort only smoke cigarettes, 0 else.
- population: the population, in hundreds of thousands, of the age group *and* smoking status cohort. Denote this P_i .
- deaths: number of lung cancer deaths in the of the specific cohort over the course of a year

. Denote this D_i .

As D_i is count data, a Poisson random variable would be appropriate to model it.

$$D_i \sim \text{Poi}(\lambda_i)$$

Of course, we don't know λ_i , but we can hypothesize that it is a function the age and smoking status. The simplest way to connect λ_i and cohort variables age and smoking statuses (not population), denoted \mathbf{x}_i , is with a link function:

$$\lambda_i = P_i \exp(\beta^T \mathbf{x}_i)$$

where β are coefficients to be determined. We require the \exp link function because the linear combination of variables may be negative, but we require the number of deaths to be positive.

Why did we put the P_i in front of the exponential? This separates the *rate of deaths*, given variables \mathbf{x}_i and represented by the exponential term, and the *number of expected deaths*, represented on the left-hand side. A larger population will naturally have more deaths, independent of the age and smoking statuses, so we need to account for this. Suppose the exponential term is 0.01, and the population is 10,000, the the expected number of deaths is $10,000 \times 0.001 = 10$. Another way to see this is show the equivalent form:

$$\frac{\lambda_i}{P_i} = \exp(\beta^T \mathbf{x}_i)$$

where the left hand side is the expected number of deaths over the cohort's population (which is a rate), and the right hand side is an estimate of the rate (given we know β_i 's). With respect to the original form, we can rewrite this as:

$$\lambda_i = \exp(\beta^T \mathbf{x}_i + \log P_i)$$

which is what more statisticians do. This is also called *adding an offset term*.

This example is quite different from our last example on text-messaging rates, though the two look similar. We are not trying to estimate a *global* λ , that is a single parameter λ that determines the distributions of all the observations, but we are actual trying to model a unique λ for each data point using the observed variables, i.e. $\lambda_i = f(\mathbf{x}_i, P_i, \beta)$. to your model.

Our conclusions are determined by the posterior distributions of β_1, β_2 and β_3 . If the distributions are shifted to be positive, then a 1 in a smoking status will shift the λ_i forward, resulting in an increase in the number of deaths. Our task in now to find the posteriors of β_i . We first need a prior for the β s: the most natural being the Normal distrubition, described below.

Normal distributions

A Normal random variable, denoted $X \sim N(\mu, 1/\tau)$, has a distribution with two parameters: the mean, μ , and the *precision*, τ . Those familiar with the Normal distribution already have probably seen σ^2 instead of τ . They are infact recipricals of each other. The change was motivated by easier

mathematics and is an artifact of Bayesian methods. Just remember: The smaller τ , the larger the spread of the distribution (i.e. we are more uncertain); the larger τ , the tighter the distribution (i.e. we are more certain). Regardless, τ is always positive.

The probability density function of a $N(\mu, 1/\tau)$ random variable is:

$$f(x|\mu, \tau) = \sqrt{\frac{\tau}{2\pi}} \exp\left(-\frac{\tau}{2}(x - \mu)^2\right)$$


We plot some different density functions below.

```
import scipy.stats as stats
nor = stats.norm
x = np.linspace( -8, 7, 150 )
mu = ( -2, 0, 3 )
tau = ( .7, 1, 2.8 )
colors = [ "#348ABD", "#A60628", "#7A68A6" ]
parameters = zip( mu, tau, colors )

for _mu, _tau, _color in parameters:
    plt.plot( x, nor.pdf( x, _mu, scale = 1./_tau ), \
              label = "$\mu = %d, \tau = %.1f$" % (_mu, _tau), color = _color )
    plt.fill_between( x, nor.pdf( x, _mu, scale = 1./_tau ), color = _color, \
                     alpha = .33)

plt.legend(loc = "upper right")
plt.xlabel("$x$")
plt.ylabel("density function at $x$")
plt.title( "Probability distribution of three different Normal random variables" )
```

<matplotlib.text.Text at 0x27d3a0b8>



A Normal random variable can be any real number, but the variable is very likely to be relatively close to μ . In fact, the expected value of a Normal is equal to its μ parameter:

$$E[X|\mu, \tau] = \mu$$

and its variance is equal to the inverse of τ :

$$\text{Var}(X|\mu, \tau) = \frac{1}{\tau}$$

Smoking and Cancer

We will model the β_i as Normal distributions. Let's start the PyMC code:

```
import pymc as mc

data = np.genfromtxt( "chp2data/smoking_death.csv", skip_header = 1,
                    delimiter=";", dtype = float )

population = data[:, -2].copy()
deaths = data[:, -1].copy()
data[:, -2] = 1 #replace the last column with a constant to represent beta_0
data = data[:, :-1]

#Instead of creating variable beta_1, beta_2, etc.,
beta = mc.MvNormal( "beta_coefs", mu = np.zeros(5) , \
                  tau = 0.0001*np.identity(5), value = np.zeros(5))
print "initial beta.value = ", beta.value

#we'll create a deterministic function that represents the exponential
#of a linear combination
@mc.deterministic
def exp_lin_comb( beta = beta ):
    return np.exp( np.dot( data, beta ) + np.log(population) )

observations = mc.Poisson( "obs", exp_lin_comb, value = deaths, observed = True )

initial beta.value = [ 0.  0.  0.  0.  0.]
```

```
model = mc.Model( [observations, beta, exp_lin_comb] )
```

```
#mysterious code to be explained in Chapter 3
```

```
map_ = mc.MAP( model )
map_.fit()
mcmc = mc.MCMC( model )
mcmc.sample( 1000000, 900000, 2 )
```

```
[*****100%*****] 1000000 of 1000000 complete
```

```
figsize(9, 10)
#histogram of the samples:
plt.figure()

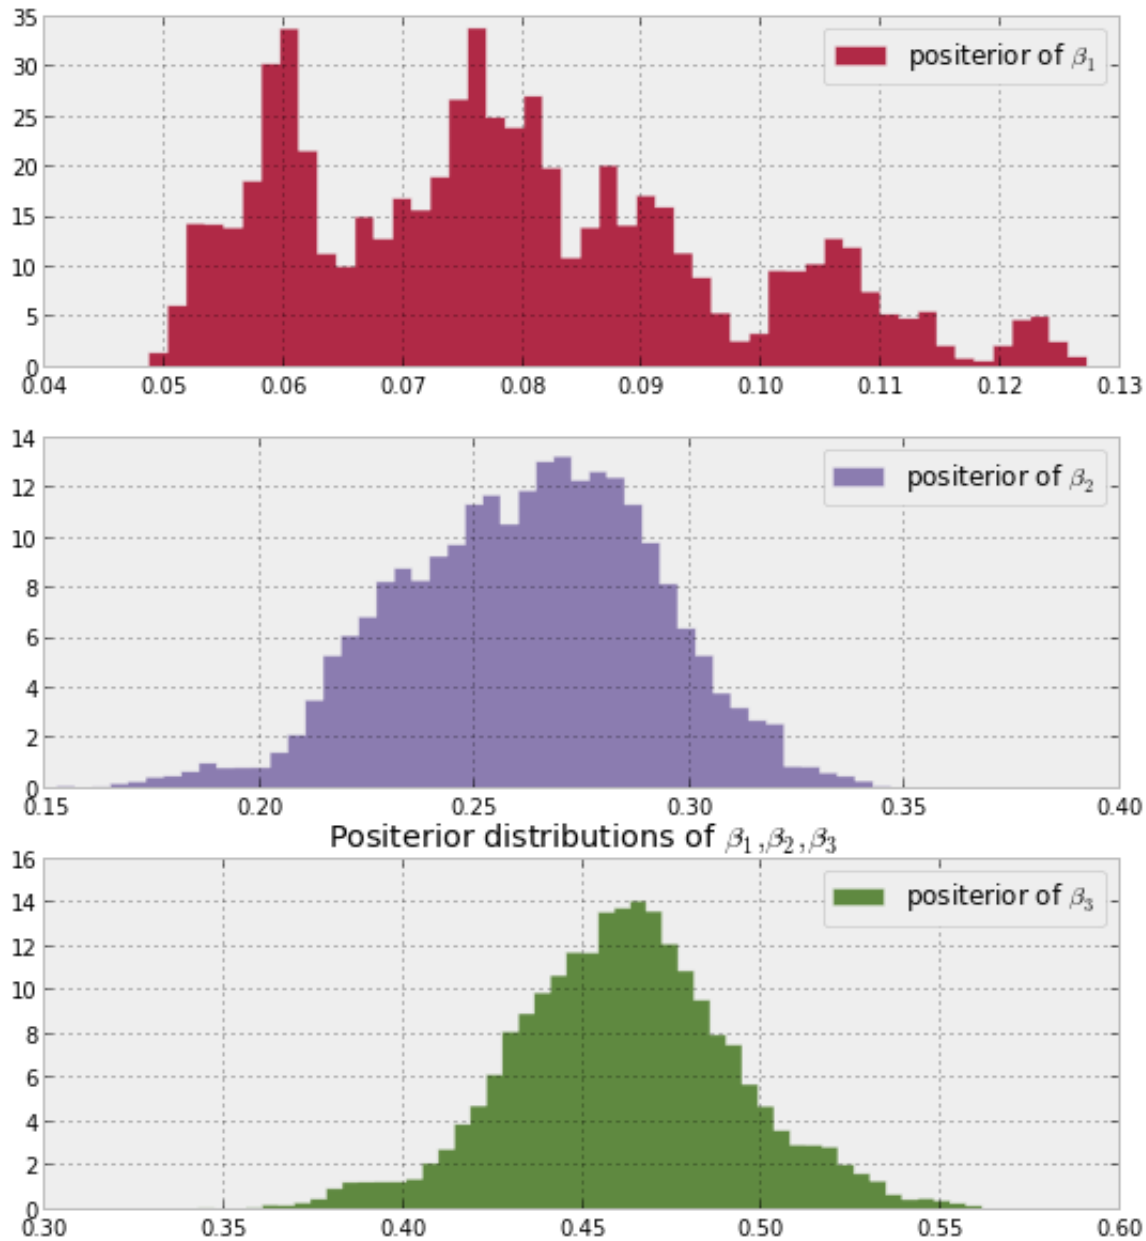
plt.subplot(311)
plt.title( r"Positerior distributions of $\beta_1, \beta_2, \beta_3$" )
plt.hist( beta_samples[:, 1], histtype='stepfilled', bins = 50, alpha = 0.85, \
        label = r"positerior of $\beta_1$", color = "#A60628", normed = True )
plt.legend()

plt.subplot(312)
plt.hist( beta_samples[:, 2], histtype='stepfilled', bins = 50, alpha = 0.85, \
        label = r"positerior of $\beta_2$", color = "#7A68A6", normed = True )
plt.legend()
```



```
plt.subplot(313)
plt.hist( beta_samples[:, 3], bins = 50, alpha = 0.85,
        label = r"positerior of $\beta_3$", \
        color="#467821", normed = True, histtype='stepfilled' )
plt.legend()
print
```

<matplotlib.text.Text at 0x193df1d0>



What can we say? It looks like all positerior distributions are strictly positive. The odd shape of the posterior of β_1 is noteworthy.

Let's perform some prediction. As we modeled the parameter λ_i in a Poisson distribution, we can find the *expected rate* of deaths. We are taking the average over the posterior distributions:

We compute `np.dot(data, beta_samples.T)`. This produces a very large matrix which is each posterior sample multiplied with the each data:

$$\beta_j^T x_i, \text{ for all } i, \text{ for all } j$$

we then exponentiate this matrix and take the mean over all j (over the posterior samples).

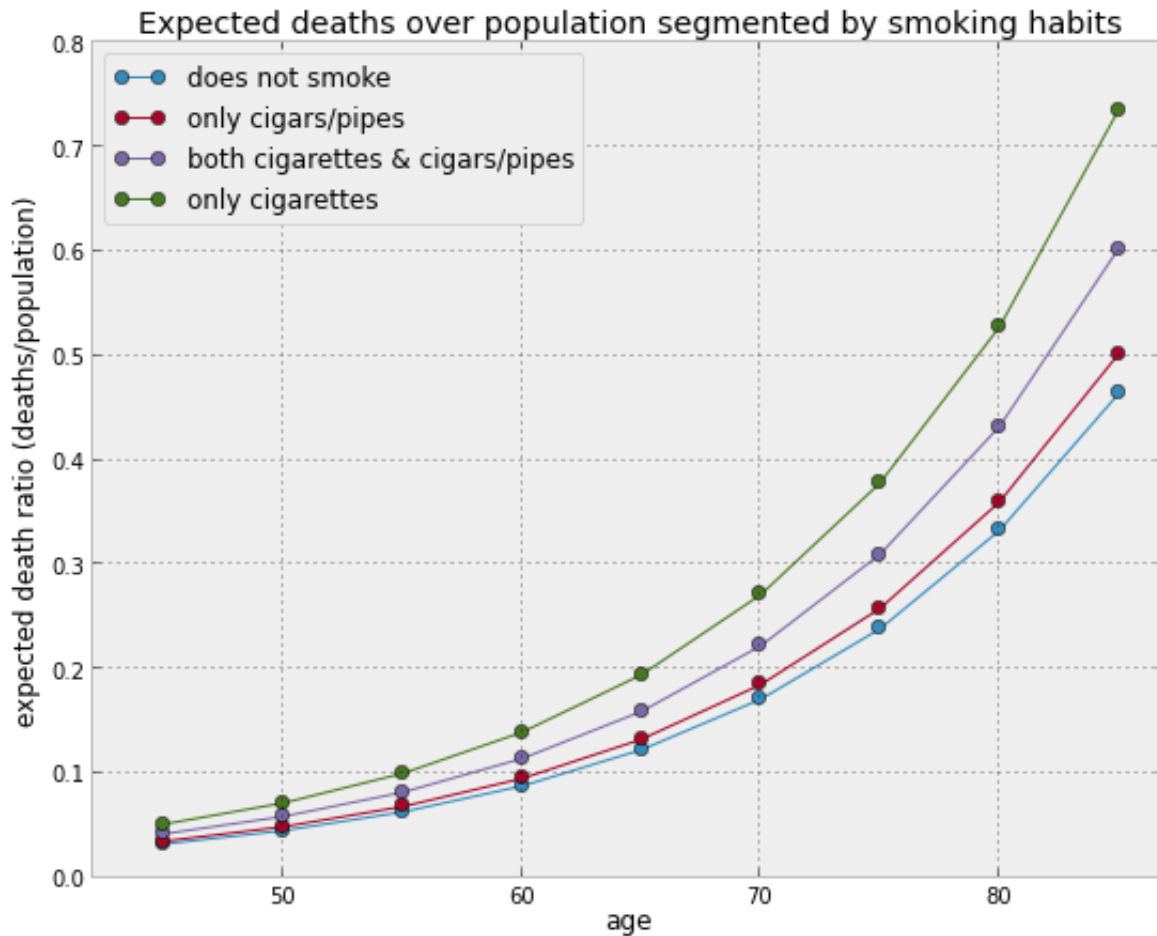
```
figsize(9, 7)
rates = np.mean( np.exp( np.dot( data, beta_samples.T ) ), axis = 1 )

colors = ["#348ABD", "#A60628", "#7A68A6", "#467821"]
labels = ["does not smoke",
          "only cigars/pipes",
          "both cigarettes & cigars/pipes",
          "only cigarettes"]

for i in range(4):
    plt.plot( 40 + 5*data[9*i:9*(i+1),0], rates[9*i:9*(i+1)], marker = 'o',
              color = colors[i], label = labels[i], lw=1)

plt.legend(loc="upper left")
plt.xlabel( "age" )
plt.ylabel("expected death ratio (deaths/population)" )
plt.title( "Expected deaths over population segmented by smoking habits")
plt.xlim( 42, 87 )
```

(42, 87)



Clearly smoking affects how long you live.

We can interpret these rates as probabilities, for example, 0.032 is the rate of a non-smokers dying

at age 40, which is equivalent to the probability of a randomly selected non-smoker aged 40-45 will die. Hence $1 - 0.032 = 0.968$ is the probability he/she survives to 45. Hence, $(1 - 0.032)(1 - 0.045)$ is the probability a non-smoker, aged 40 will survive until 50, etc. Let's plot this.


```
figsize( 9, 3.5 )
prob_survival = 1-rates

colors = ["#348ABD", "#A60628", "#7A68A6", "#467821"]
labels = ["does not smoke",
          "only cigars/pipes",
          "both cigarettes & cigars/pipes",
          "only cigarettes"]

for i in range(4):
    _p = prob_survival[ 9*i:9*(i+1) ].cumprod()
    plt.plot( 40 + 5*data[9*i:9*(i+1),0], _p, marker = 'o',
             color = colors[i], label = labels[i], lw=1)

plt.legend(loc="lower left")
plt.xlabel("Age")
plt.ylabel("Probability of survival")
plt.title("Probability of survival, segmented by smoking")
```

<matplotlib.text.Text at 0x21a350f0>



It appears that smoking cigarettes will almost certainly kill you before age 85, whereas not smoking provides you with 18% chances to live this long. So, for you young folks out there, which path do you want to take?

Example: Challenger Space Shuttle Disaster

On January 28, 1986, the twenty-fifth flight of the U.S. space shuttle program ended in disaster when one of the rocket boosters of the Shuttle Challenger exploded shortly after lift-off, killing all seven crew members. The presidential commission on the accident concluded that it was caused by the failure of an O-ring in a field joint on the rocket booster, and that this failure was due to a faulty design that made the O-ring unacceptably sensitive to a number of factors including outside

temperature. Of the previous 24 flights, data were available on failures of O-rings on 23, (one was lost at sea), and these data were discussed on the evening preceding the Challenger launch, but unfortunately only the data corresponding to the 7 flights on which there was a damage incident were considered important and these were thought to show no obvious trend. The data are shown below (see [1]):


```
figsize( 9, 3.5 )
challenger_data = np.genfromtxt("chp2data/challenger_data.csv", skip_header = 1, \
                                usecols=[1,2], missing_values="NA", delimiter=",")
#drop the NA values
challenger_data = challenger_data[ ~np.isnan(challenger_data[:,1]) ]


#plot it, as a function of tempature (the first column)
print "Temperature, O-Ring failure"
print challenger_data

plt.scatter( challenger_data[:,0], challenger_data[:,1], s = 75, color="k",
            alpha = 0.5)
plt.yticks([0,1])
plt.ylabel("Damage Incident?")
plt.xlabel("Outside temperature (Farhenhit)" )
plt.title("Defects of the Space Shuttle O-Rings vs temperature")
print
```

Temperature, O-Ring failure

```
[[ 66.  0.]
 [ 70.  1.]
 [ 69.  0.]
 [ 68.  0.]
 [ 67.  0.]
 [ 72.  0.]
 [ 73.  0.]
 [ 70.  0.]
 [ 57.  1.]
 [ 63.  1.]
 [ 70.  1.]
 [ 78.  0.]
 [ 67.  0.]
 [ 53.  1.]
 [ 67.  0.]
 [ 75.  0.]
 [ 70.  0.]
 [ 81.  0.]
 [ 76.  0.]
 [ 79.  0.]
 [ 75.  1.]
 [ 76.  0.]
 [ 58.  1.]]
```





It looks clear that *the probability* of damage incidents occuring increases as the outside temperature decreases. We are interested in modeling the probability here because it does not look like there is a strict threshold between temperature and a damage incident.


We need a function of temperature, call it $p(t)$, that is bounded between 0 and 1 (so as to model a probability) and changes from 0 to 1 as we decrease temperature. There are actually many such functions, but the most popular choice is the *logistic function*.

$$p(t) = \frac{1}{1 + e^{\beta t}}$$

In this model, β is the variable we are uncertain about. Below is the function plotted for $\beta = 1, 3, -5$.

```
def logistic( x, beta):
    return 1.0/( 1.0 + np.exp( beta*x) )
x = np.linspace( -4, 4, 100 )
plt.plot(x, logistic( x, 1), label = r"$\beta = 1$")
plt.plot(x, logistic( x, 3), label = r"$\beta = 3$")
plt.plot(x, logistic( x, -5), label = r"$\beta = -5$")
plt.legend()
```

<matplotlib.legend.Legend at 0x2226cda0>




But something is missing. In the plot above, the probability changes quickly only near zero, but in our data above the probability changes around 65 to 70. We need to add a *bias* term to our logistic function:

$$p(t) = \frac{1}{1 + e^{\beta t + \alpha}}$$

This means we can change where the curve changes. Some plots are below, with differing α .

```
def logistic( x, beta, alpha):
    return 1.0/( 1.0 + np.exp( beta*x + alpha) )
x = np.linspace( -4, 4, 100 )
plt.plot(x, logistic( x, 1, 1), label = r"$\beta = 1, \alpha = 1$")
plt.plot(x, logistic( x, 3, -2), label = r"$\beta = 3, \alpha = -2$")
plt.plot(x, logistic( x, -5, 7), label = r"$\beta = -5, \alpha = 7$")
plt.legend(loc="lower left")
```

<matplotlib.legend.Legend at 0x29ffe588>



Adding a constant term α amounts to shifting the curve left or right (hence why it called a *bias*.)

Let's start modeling this in PyMC. The β, α paramters have no reason to be positive, bounded or relatively large, so there are best modeled by a Normal random variable.

```
import pymc as mc

temperature = challenger_data[:,0]
defect_ind= challenger_data[:,1]

beta = mc.Normal( "beta", 0, 0.001, value = 0 )
alpha = mc.Normal( "alpha", 0, 0.001, value = 0 )

@mc.deterministic
def p( temp = temperature, alpha = alpha, beta = beta):
    return 1.0/( 1. + np.exp( beta*temperature + alpha) )
```

We have our probabilities, but how do we connect them to our observed data? A *Bernoulli* random variable with parameter p (denoted $\text{Ber}(p)$), is a random variable that takes value 1 with probability p , and 0 else. Thus, our model can look like:

$$\text{Defect Incident, } D_i \sim \text{Ber}(p(t_i)), \quad i = 1..N$$

where $p(t)$ is our logistic function and t_i are the temperatures we have observations about. Notice in the above code we had to set the values of `beta` and `alpha` to 0. The reason for this is that if `beta` and `alpha` are very large, they make `p` equal to 1 or 0. Unfortunately, `mc.Bernoulli` does not like probabilities of 0 or 1, though they are mathematically well-defined probabilities. So by

setting the values to 0, we set p to a reasonable starting value. This has no effect on our results, nor does it mean we are including information in our prior. It is simply a computational caveat in PyMC.

```
p.value
```

```
array([ 0.5,  0.5,  0.5,  0.5,  0.5,  0.5,  0.5,  0.5,  0.5,  0.5,  0.5,
        0.5,  0.5,  0.5,  0.5,  0.5,  0.5,  0.5,  0.5,  0.5,  0.5,
        0.5])
```

```
observed = mc.Bernoulli( "bernoulli_obs", p, \
                        value = defect_ind, observed=True)

model = mc.Model( {"observed":observed, "beta":beta, "alpha":alpha} )

#mysterious code to be explained in Chapter 3
map_ = mc.MAP(model)
map_.fit()
mcmc = mc.MCMC( model )
mcmc.sample( 260000, 220000, 3 )
```

```
[*****100%*****] 260000 of 260000 complete
```

We have trained our model on the observed data, now we can sample values from the posterior.
Let's look at the posterior distributions for α and β :


```
alpha_samples = mcmc.trace( 'alpha' )[:]
beta_samples = mcmc.trace( 'beta' )[:]

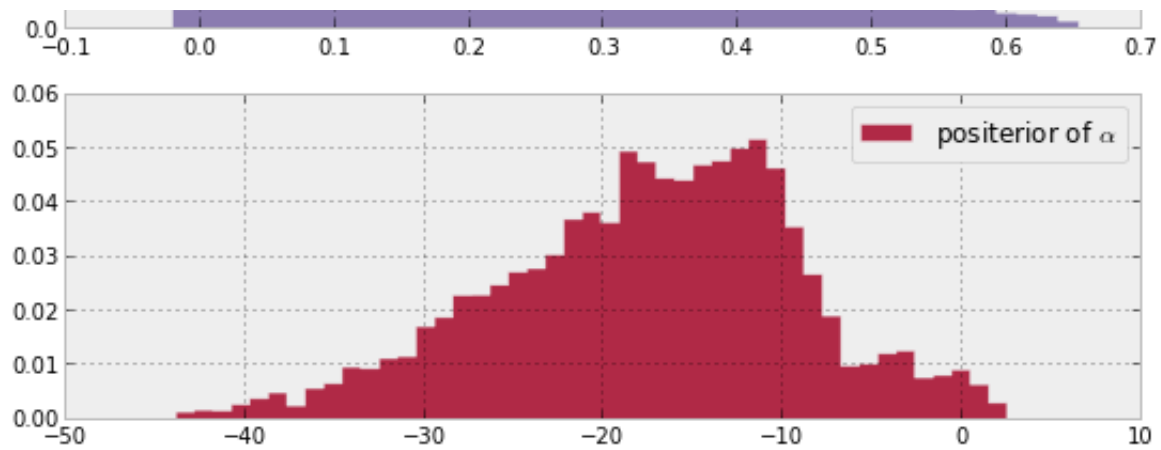
figsize(9, 6)

#histogram of the samples:
plt.subplot(211)
plt.title(r"Posterior distributions of the variables  $\alpha, \beta$ ")
plt.hist( beta_samples, histtype='stepfilled', bins = 45, alpha = 0.85, \
        label = r"posterior of  $\beta$ ", color = "#7A68A6",normed = True )
plt.legend()

plt.subplot(212)
plt.hist( alpha_samples, histtype='stepfilled', bins = 45, alpha = 0.85, \
        label = r"posterior of  $\alpha$ ", color = "#A60628",normed = True )
plt.legend()
```

```
<matplotlib.legend.Legend at 0x2a94d0b8>
```





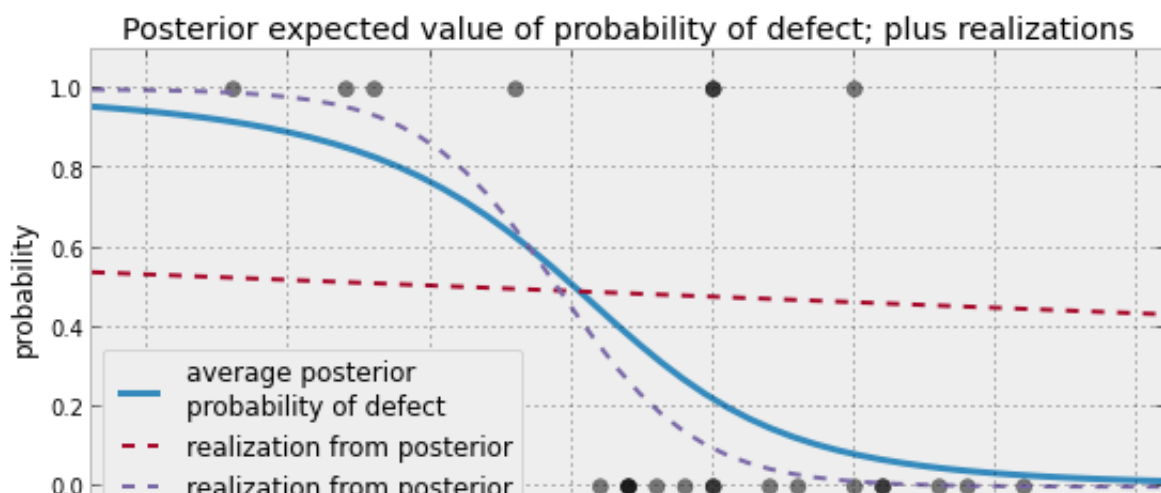
Regarding the spread of the data, we are very uncertain about what the parameters may be (though considering the low sample size and the large overlap of defects to no-defects this behaviour is perhaps expected). The alternative, frequentist-logistic regression would return two scalars that are likely very different from the true values.


Next, let's look at the *expected probability* for a specific value of the temperature.

```
t = np.linspace( temperature.min() - 5, temperature.max()+5, 50 )
linear_combination= np.dot( beta_samples[:,None], t[None,:] ) + alpha_samples[:,None]
p = 1.0/(1.0 + np.exp( linear_combination ) )
mean_prob_t = p.mean(axis=0)
```

```
figsize( 9, 4)
plt.plot( t, mean_prob_t, lw = 3, label = "average posterior \nprobability \n of defect")
plt.plot( t, p[0, :], ls="--",label="realization from posterior" )
plt.plot( t, p[-2, :], ls="--", label="realization from posterior" )
plt.scatter( temperature, defect_ind, color = "k", s = 50, alpha = 0.5 )
plt.title("Posterior expected value of probability of defect; plus realizations")
plt.legend(loc= "lower left")
plt.ylim( -0.1, 1.1 )
plt.xlim( t.min(), t.max() )
plt.ylabel("probability")
plt.xlabel("temperature")
```

<matplotlib.text.Text at 0x25d4ff28>






Above we also plotted two possible realizations of what the actual underlying system might be. Both are as equally likely as any other. Notice they end to deviate from the expected value line around 60 degrees. An interesting question is for what temperatures are we most uncertain about the probability being? Below we plot the expected value line and the associated distribution at each temperature.

```
plt.plot( t, mean_prob_t, lw = 3, label = "average posterior \nprobability \nof defect")
n = 2000
for i, _t in enumerate(t[:-1]):
    plt.scatter( _t*ones( n ) + 0.20*np.random.randn(n), p[:n,i], alpha = 0.009,\
                color = "#D55E00")

#what's a better way to only have a for loop and include it in the legend? and
# have it somewhat show. You can see an artifact at t = 85.
plt.scatter( t[-1]*ones( n ) + 0.20*np.random.randn(n), p[:n,-1], alpha = 0.3,\
            color = "#D55E00", label = "samples from posterior at temp $t$")

plt.xlim( t.min(), t.max() )
plt.legend()
plt.scatter( temperature, defect_ind, color = "k", s = 50, alpha = 0.5 )
plt.xlabel("temp, $t$")
plt.ylabel("probability estimate" )
```

<matplotlib.text.Text at 0x25aef278>



We can see that as the temperature nears 60 degrees, the distributions spread out over $[0,1]$ quickly. As we pass 70 degrees, the distributions tighten again. This can give us insight about how to proceed next: we should probably test more O-rings around 60-65 temperature to get a better estimate of probabilities in that range.


What about the day of the Challenger disaster?

On the day of the Challenger disaster, the outside temperature was 31 degrees fahrenheit. What is the posterior distribution of a defect, given this temperature? The distribution is plotted below. It looks almost guaranteed that the Challenger was going to be subject to defective O-rings.

```
figsize(9, 2.5)
t = 31
p = beta_samples*31 + alpha_samples
prob_31 = 1.0/(1.0 + exp(p) )

plt.xlim( 0.99, 1)
plt.hist( prob_31, bins =160, normed = True, histtype='stepfilled' )
plt.title( "Posterior distribution of probability of defect, given $t = 31$" )
plt.xlabel( "probability of defect occurring in O-ring" )

print
```




Is there really a relationship between failure and temperature?

An criticism of our above analysis is that *assumed* that the relationship followed a logistic model, this we implicitly assumed that the probabilities change over temperature. Let's look at the data again:

```
plt.scatter( challenger_data[:,0], challenger_data[:,1], s = 75, color="k",
            alpha = 0.5)
plt.yticks([0,1])
plt.ylabel("Damage Incident?")
plt.xlabel("Outside temperature (Farhenhit)" )
plt.title("Defects of the Space Shuttle O-Rings vs temperature")
```

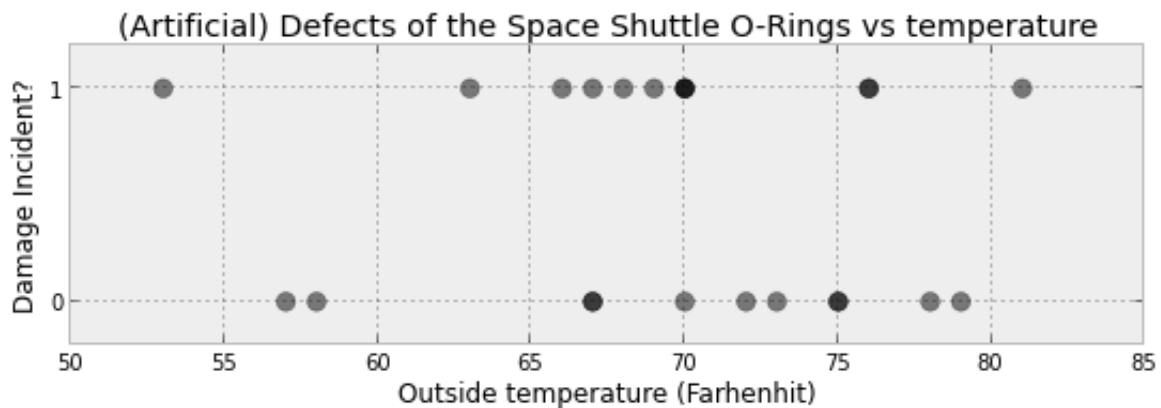
<matplotlib.text.Text at 0x28f89c18>



50 55 60 65 70 75 80 85
Outside temperature (Farhenhit)

Could it be that infact this pattern occured by chance? This might explain the large overlap in temperatures. After all, I can produce similar plots:

```
figsize( 9, 2.5)
n = challenger_data.shape[0]
plt.scatter( challenger_data[:,0], stats.bernoulli.rvs(0.6, size = n) , s = 75, color
alpha = 0.5)
plt.yticks([0,1])
plt.ylabel("Damage Incident?")
plt.xlabel("Outside temperature (Farhenhit)" )
plt.title("(Artificial) Defects of the Space Shuttle O-Rings vs temperature")
print
```



Introducing the Bayes factor

The *Bayes factor* is a measure comparing two models. In our example, on one hand we believe that temperature plays an important role is the probability of O-ring failures. On the other hand, we believe that there is no connection, and the appears like that by chance. We can compare these model using the ratio of the *probabilities of observing the data, given the model*:

$$\frac{P(\text{observations}|\text{Model}_1)}{P(\text{observations}|\text{Model}_2)}$$

Calculating this is not at all obvious. For now, let's select a set of parameters from the posterior distribution.

```
alpha_hat = alpha_samples[0] #select the first alpha samples
beta_hat = beta_samples[0] #select the first beta sample

observed_temperatures = challenger_data[:,0]
D = challenger_data[:,1] #defect or not?

p_hat = 1.0/(1.0 + np.exp(beta_hat*challenger_data[:,0] + alpha_hat ) )
print "estimates of probability at observed temperature, defects: "
print np.array(zip(p_hat, observed_temperatures, D) )[:3, :]
```

estimates of probability at observed temperature, defects:

```
[[ 0. 66. 0.]  
 [ 0. 70. 1.]  
 [ 0. 69. 0.]
```

Assuming this model, that is with set $\hat{\beta}, \hat{\alpha}$, the probability of our observations is equal to the product of:

$$P(\text{Ber}(p(t_i | \hat{\beta}, \hat{\alpha})) = D_i), \quad i = 1..N$$

For example, using the output above

$$\begin{aligned} \text{Ber}(p(t_1 | \hat{\beta}, \hat{\alpha})) &= D_1 \\ \Rightarrow \text{Ber}(0.667) &= 0 \end{aligned}$$

And the probability of this occurring is $(1 - 0.667) = 0.333$ (Recall the definition of a Bernoulli random variable $\text{Ber}(p)$: it is equal to 1 with probability p .) As each observation is independent, we multiply all the observation together. A clever way to do this is, for a specific i :

$$(p(t_1 | \hat{\beta}, \hat{\alpha}))^{D_i} (1 - p(t_1 | \hat{\beta}, \hat{\alpha}))^{(1-D_i)}$$

(For many observations, this quantity can overflow, so is recommended to take the log of the above::

$$D_i \log(p(t_1 | \hat{\beta}, \hat{\alpha})) + (1 - D_i) \log(1 - p(t_1 | \hat{\beta}, \hat{\alpha}))$$

and sum the terms instead of multiplying. Be sure use this transform for both models you are comparing.)

```
product = p_hat**( D )*(1-p_hat)**(1-D)  
p_model_1 = product.prod()  
print p_model_1
```

2.84378201156e-05

The second model we are comparing against is that $\beta = 0$, i.e. there is no relationship between probabilities and temperature. We perform the same calculations as above. Notice that `beta=0` here in the below code:

```
beta = 0  
alpha = mc.Normal( "alpha", 0, 0.001, value = 0 )  
  
@mc.deterministic  
def p( temp = temperature, alpha = alpha, beta = beta):  
    return 1.0/( 1. + np.exp( beta*temperature + alpha) )  
  
observed = mc.Bernoulli( "bernoulli_obs", p, \  
                        value = defect_ind, observed=True)  
  
model = mc.Model( {"observed":observed, "beta":beta, "alpha":alpha} )  
  
#mysterious code to be explained in Chapter 3  
map_ = mc.MAP(model)
```

```

map_.fit()
mcmc = mc.MCMC( model )
mcmc.sample( 260000, 220000, 3 )

alpha_samples_model_2 = mcmc.trace("alpha")[:]
alpha_hat = alpha_samples_model_2[0]
beta_hat = 0

p_hat = 1.0/(1.0 + np.exp(beta_hat*challenger_data[:,0] + alpha_hat ) )
print "estimates of probability at observed temperature, defects: "
print np.array(zip(p_hat, observed_temperatures, D )[:3, :])

[*****100%*****] 260000 of 260000 complete
estimates of probability at observed temperature, defects:
[[ 0.323  66.      0.   ]
 [ 0.323  70.      1.   ]
 [ 0.323  69.      0.   ]]

```

```

#compute the probability of observations given the model_2

```

```

product = p_hat** ( D )*(1-p_hat)**(1-D)
model_2 = product.prod()
print model_2

```

```

4.28136759838e-07

```

```

print "Bayes factor = %.3f"%(model_1/model_2)

```

```

Bayes factor = 66.422

```

Is this good? Below is a chart that can be used to compare the computed Bayes factors to a degree of confidence in M1.

- <1:1 Negative (supports M2)
- 1:1 to 3:1 Barely worth mentioning
- 3:1 to 10:1 Substantial
- 10:1 to 30:1 Strong
- 30:1 to 100:1 Very strong
- > 100:1 Decisive

We are not done yet. If you recall, we selected only one model, but recall we actually have a distribution of possible models (sequential pairs of (β, α) from the posterior distributions). So to be correct we need to average over samples from the posterior:

```

linear_combination= np.dot( beta_samples[:,None], observed_temperatures[None,:]) \
                    + alpha_samples[:,None]

p_model_1 = 1.0/(1.0 + np.exp( linear_combination ) )
product = p_model_1** ( D )*(1-p_model_1)**(1-D)
model_1_product = product.prod(axis=1).mean()
print model_1_product

p_model_2 = 1.0/(1.0 + np.exp( alpha_samples_model_2 ) )[:,None]
product = p_model_2** ( D )*(1-p_model_2)**(1-D)

```

```

model_2_product = product.prod(axis=1).mean()
print model_2_product

print

print "Bayes factor: %.3f"%(model_1_product/model_2_product)

```

1.83483099109e-05

5.06760508661e-07

Bayes factor: 36.207

It looks we have a pretty good model, and temperature *is* significant.

Exercises

1. How would add a prior belief that smoking causes more deaths in first example?
2. Try plotting α samples versus β samples.

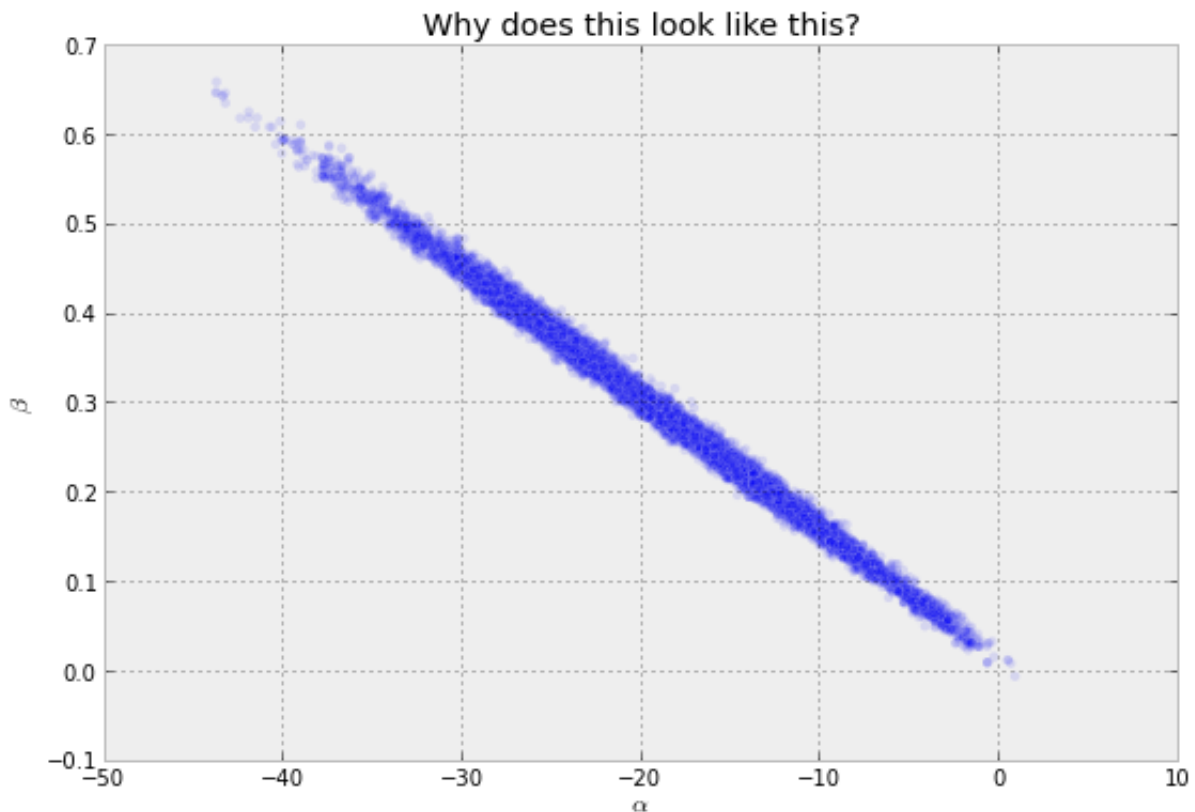
#type your code here.

```

plt.scatter( alpha_samples, beta_samples, alpha = 0.1 )
plt.title( "Why does the plot look like this?" )
plt.xlabel( r"$\alpha$" )
plt.ylabel( r"$\beta$" )

```

<matplotlib.text.Text at 0x25a65438>



References

- [1] Dalal, Fowlkes and Hoadley (1989), JASA, 84, 945-957.
- [2] German Rodriguez. Datasets. In WWS509. Retrieved 30/01/2013, from <http://data.princeton.edu/wws509/datasets/#smoking>.